

A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

Machine Learning Project

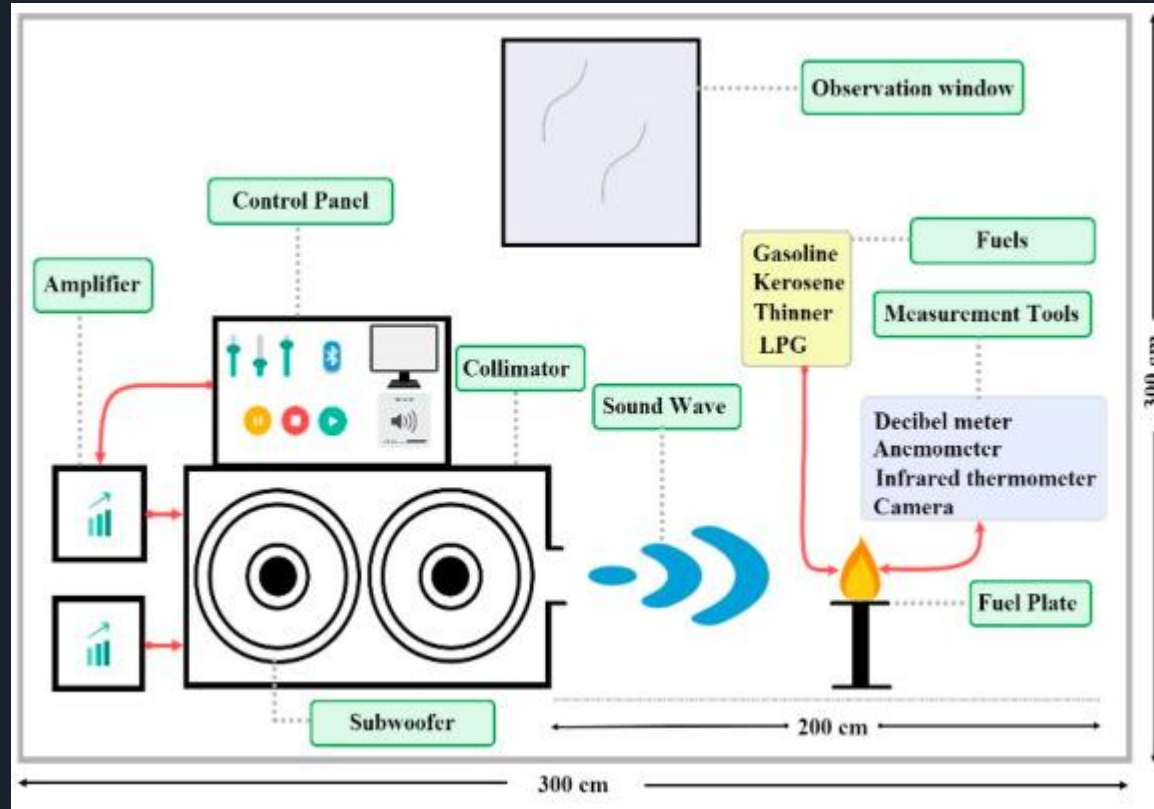
Thomas D and Josh Q



Data Source and Background

- Data is taken from Kaggle:
<https://www.kaggle.com/datasets/muratkokludataset/acoustic-extinguisher-fire-dataset?resource=download>
- The data is from research done on sound waves from speakers putting out fires
- 7 columns and 17,442 rows of data

Study Graphic





Data Overview

Size	Fuel	Distance	Decibel	Airflow	Frequency	Status
1	gasoline	10	96	0	75	0
1	gasoline	10	96	0	72	1
1	gasoline	10	96	2.6	70	1


- Fuel: type of fuel used to start fire
- Size: size of canister used to hold fuel
- Distance: Distance of fire from speaker
- Decibel: Volume of speaker
- Airflow: Airflow produced by the speaker
- Frequency: Frequency of sound from the speaker
- Status: Whether or not the fire was put out

Data Overview Cont.

SIZE	FUEL	DISTANCE	DESIBEL
Min. :1.000	Length:17442	Min. : 10	Min. : 72.00
1st Qu.:2.000	Class :character	1st Qu.: 50	1st Qu.: 90.00
Median :3.000	Mode :character	Median :100	Median : 95.00
Mean :3.412		Mean :100	Mean : 96.38
3rd Qu.:5.000		3rd Qu.:150	3rd Qu.:104.00
Max. :7.000		Max. :190	Max. :113.00
AIRFLOW	FREQUENCY	STATUS	
Min. : 0.000	Min. : 1.00	Min. :0.0000	
1st Qu.: 3.200	1st Qu.:14.00	1st Qu.:0.0000	
Median : 5.800	Median :27.50	Median :0.0000	
Mean : 6.976	Mean :31.61	Mean :0.4978	
3rd Qu.:11.200	3rd Qu.:47.00	3rd Qu.:1.0000	
Max. :17.000	Max. :75.00	Max. :1.0000	

- Nearly equal amount of fires put out and fires not put out
- Equal distribution of each size for each fuel type

Correlation Plot





Data Preprocess

- Converted the three fuels into an integers.
- Normalized numeric data using min-max scale to get numbers on a similar scale
- Converted Fuel into categorical data (gasoline, thinner, kerosene, liquid petroleum gas turn into 0,1,2,3).
- Measurement of size was different for LPG than other fuels, so all LPG entries removed from the data set

Size	Fuel	Distance	Decibel	Airflow	Frequency	Status
1	0	10	0.05263158	0.0000000	1.0000000	0
1	0	10	0.05263158	0.0000000	0.9600000	1
1	0	10	0.05263158	0.1529412	0.9333333	1



Training and Testing Data

- Training data was 12,000 randomly selected data points. The testing was then the 3,390 remaining data points
- Seed was set to make the testing and training results reproducible



Neural Network Architecture

- 3 layers
- 64 nodes for the two ReLU layers
- 1 node for the output layer, it uses sigmoid
- Loss: binary cross entropy
- Optimizer: adam
- 35 epochs
- Training loss: loss: 0.1119
- Training accuracy: 0.9503
- Testing loss: 0.1438
- Testing accuracy: 0.9342

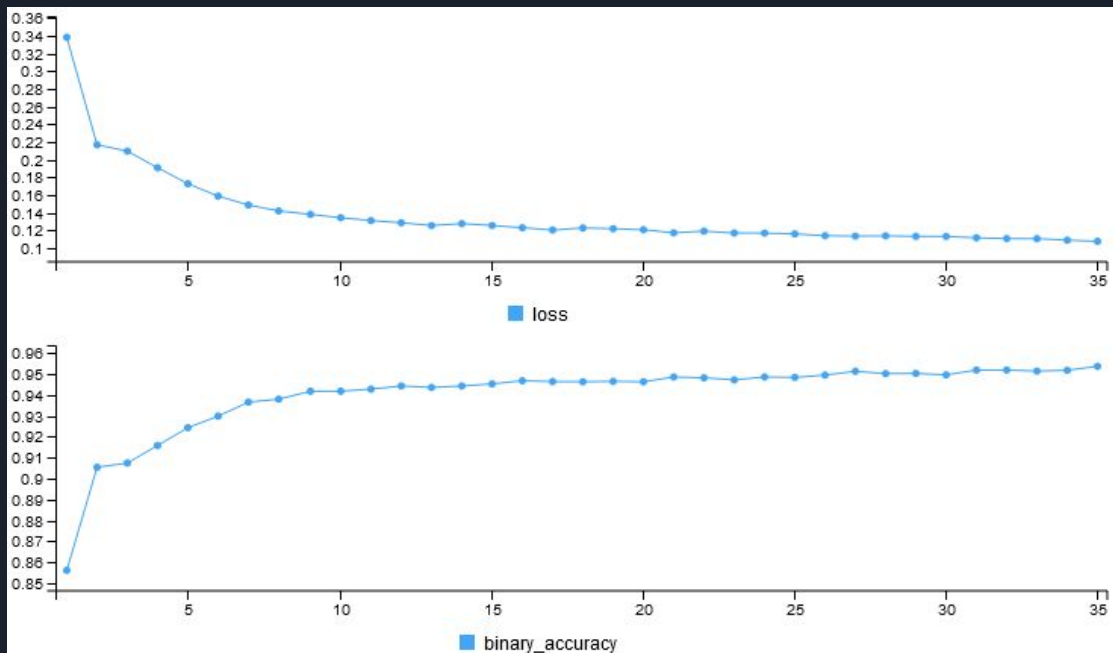
```
#64 nodes
first_model = keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")

first_model %>% compile(
  loss = "binary_crossentropy",
  metrics = list("binary_accuracy"),
  optimizer = "adam"
)

first_history <- first_model %>%
  fit(
    x = xTr,
    y = yTr, # label is the last column
    epochs = 35,
    validation_data = list(xTest, yTest),
    verbose = 2
  )
```

Run of the model

```
Epoch 1/35  
375/375 - 2s - loss: 0.3384 - binary_accuracy: 0.8562 - 2s/epoch - 4ms/step  
Epoch 2/35  
375/375 - 1s - loss: 0.2167 - binary_accuracy: 0.9054 - 747ms/epoch - 2ms/step  
Epoch 3/35  
375/375 - 1s - loss: 0.2095 - binary_accuracy: 0.9074 - 740ms/epoch - 2ms/step  
Epoch 4/35  
375/375 - 1s - loss: 0.1907 - binary_accuracy: 0.9158 - 721ms/epoch - 2ms/step  
Epoch 5/35  
375/375 - 1s - loss: 0.1726 - binary_accuracy: 0.9243 - 738ms/epoch - 2ms/step  
Epoch 6/35  
375/375 - 1s - loss: 0.1586 - binary_accuracy: 0.9298 - 762ms/epoch - 2ms/step  
Epoch 7/35  
375/375 - 1s - loss: 0.1485 - binary_accuracy: 0.9365 - 732ms/epoch - 2ms/step  
Epoch 8/35  
375/375 - 1s - loss: 0.1418 - binary_accuracy: 0.9379 - 726ms/epoch - 2ms/step  
Epoch 9/35  
375/375 - 1s - loss: 0.1379 - binary_accuracy: 0.9417 - 764ms/epoch - 2ms/step  
Epoch 10/35  
375/375 - 1s - loss: 0.1341 - binary_accuracy: 0.9417 - 721ms/epoch - 2ms/step  
Epoch 11/35  
375/375 - 1s - loss: 0.1309 - binary_accuracy: 0.9427 - 745ms/epoch - 2ms/step  
Epoch 12/35  
375/375 - 1s - loss: 0.1284 - binary_accuracy: 0.9442 - 740ms/epoch - 2ms/step  
Epoch 13/35  
375/375 - 1s - loss: 0.1253 - binary_accuracy: 0.9435 - 769ms/epoch - 2ms/step  
Epoch 14/35  
375/375 - 1s - loss: 0.1274 - binary_accuracy: 0.9442 - 781ms/epoch - 2ms/step  
Epoch 15/35  
375/375 - 1s - loss: 0.1253 - binary_accuracy: 0.9452 - 729ms/epoch - 2ms/step  
Epoch 16/35  
375/375 - 1s - loss: 0.1228 - binary_accuracy: 0.9467 - 729ms/epoch - 2ms/step  
Epoch 17/35  
375/375 - 1s - loss: 0.1202 - binary_accuracy: 0.9463 - 721ms/epoch - 2ms/step  
Epoch 18/35  
375/375 - 1s - loss: 0.1225 - binary_accuracy: 0.9462 - 722ms/epoch - 2ms/step  
Epoch 19/35  
375/375 - 1s - loss: 0.1217 - binary_accuracy: 0.9464 - 762ms/epoch - 2ms/step  
Epoch 20/35  
375/375 - 1s - loss: 0.1205 - binary_accuracy: 0.9462 - 730ms/epoch - 2ms/step  
Epoch 21/35  
375/375 - 1s - loss: 0.1170 - binary_accuracy: 0.9485 - 727ms/epoch - 2ms/step  
Epoch 22/35  
375/375 - 1s - loss: 0.1189 - binary_accuracy: 0.9481 - 744ms/epoch - 2ms/step  
Epoch 23/35  
375/375 - 1s - loss: 0.1168 - binary_accuracy: 0.9471 - 717ms/epoch - 2ms/step
```





Dropout Model

- Example of 3 of the 4 hidden layers having dropout applied to them
- Training loss: 0.0942
- Training accuracy: 0.9603
- Testing loss: 0.0779
- Testing accuracy: 0.9678

```
dropout_model = keras_model_sequential() %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dropout(0.1) %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dropout(0.1) %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dropout(0.1) %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dropout(0.1) %>%  
  layer_dense(units = 1, activation = "sigmoid")  
  
dropout_model %>% compile(  
  loss = "binary_crossentropy",  
  metrics = list("binary_accuracy"),  
  optimizer = "adam"  
)  
  
dropout_history <- dropout_model %>%  
  fit(  
    x = xTr,  
    y = yTr,  
    epochs = 35,  
    validation_data = list(xTest, yTest),  
    verbose = 2  
  )
```

L2 Regularization Model

- Our model that uses L2 regularization on all 4 of the hidden layers
- Training loss: 0.1405
- Training accuracy: 0.9514
- Testing loss: 0.1300
- Testing accuracy: 0.9581

```
l2_model = keras_model_sequential() %>%
  layer_dense(units = 128, activation = "relu",
    kernel_regularizer = regularizer_l2(l1 = 0.001)) %>%
  layer_dense(units = 128, activation = "relu",
    kernel_regularizer = regularizer_l2(l1 = 0.001)) %>%
  layer_dense(units = 128, activation = "relu",
    kernel_regularizer = regularizer_l2(l1 = 0.001)) %>%
  layer_dense(units = 128, activation = "relu",
    kernel_regularizer = regularizer_l2(l1 = 0.001)) %>%
  layer_dense(units = 1, activation = "sigmoid")

l2_model %>% compile(
  loss = "binary_crossentropy",
  metrics = list("binary_accuracy"),
  optimizer = "adam"
)

l2_history <- l2_model %>%
  fit(
    x = xTr,
    y = yTr,
    epochs = 35,
    validation_data = list(xTest, yTest),
    verbose = 2
  )
```

Some tests

- 4 hidden layers, 64 nodes each, 35 epochs
 - loss: 0.0916 - accuracy: 0.9575
- 4 hidden layers, 64 nodes, dropout(0.2), 35 epochs
 - loss: 0.0925 - accuracy: 0.9631
- 4 hidden layers, 128 nodes, L2, 35 epochs
 - loss: 0.1241 - accuracy: 0.9631
- **4 hidden layers, 128 nodes each, 50 epochs**
 - **loss: 0.0711 - accuracy: 0.9714**
- 4 hidden layers, 128 nodes, dropout(0.1), 50 epochs
 - loss: 0.0775 - accuracy: 0.9673
- 6 hidden layers, 128 nodes, L2, 50 epochs
 - loss: 0.1149 - accuracy: 0.9649

```
first_model = keras_model_sequential() %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dense(units = 128, activation = "relu") %>%  
    layer_dense(units = 128, activation = "relu") %>%  
  layer_dense(units = 128, activation = "relu") %>%  
  layer_dense(units = 1, activation = "sigmoid")  
  
first_model %>% compile(  
  loss = "binary_crossentropy",  
  metrics = list("binary_accuracy"),  
  optimizer = "adam"  
)  
  
first_history <- first_model %>%  
  fit(  
    x = xTr,  
    y = yTr, # label is the last column  
    epochs = 50,  
    verbose = 2  
  )
```



Final Remarks

- Removing FUEL column of data decreased accuracy about 2%
- Normalizing the data increased accuracy about 1%
- Removing any other column reduced accuracy below 90% for all methods
- Were able to increase testing accuracy from 0.9342 to 0.9714 by adding 2 more layers and increasing the nodes and epochs
- Dropout worked better than L2 regularization for our data